

Title: Comparative Performance Analysis of Python, C++, Java, and Go

1. Objective

To benchmark and evaluate the performance and efficiency of Python, C++, Java, and Go across a common computational task using key metrics: execution time, memory usage, CPU load, and debugging simplicity.

2. Benchmark Task Chosen

We will use Fibonacci sequence calculation (recursive and iterative) and a file-processing task (reading, sorting, and writing large datasets) as representative compute-bound and I/O-bound tasks. We have also used specific language profiler for in-depth analysis.

3. Benchmark Criteria

Metric	Description
Execution Speed	Time taken to complete the task
Memory Usage	Peak RAM consumed during execution
CPU Utilization	% of CPU consumed during peak load
Debugging Ease	Developer experience and tooling
Development Speed	Time/effort to write/maintain code
Stability Under Load	Handles stress tests (large files, recursion)

Language Profiler Tool(s)

Python	cProfile, memory_profiler
C++	Visual Studio Profiler (Call Tree, CPU stats), time
Java	VisualVM, JMH
Go	Built-in pprof

5. Benchmark Code Snippets

Task 1: Fibonacci (Recursive)

Python

```
fibonacci.py X
fibonacci.py > ...
1   import time
2   def fib(n):
3       return 1 if n <= 2 else fib(n-1) + fib(n-2)
4
5   start = time.time()
6   print(fib(30))
7   print("Time:", time.time() - start)
8
```

Results

```
PROBLEMS    OUTPUT    DEBUG CONSOLE    TERMINAL    PORTS
PS C:\Assignment1> & 'c:\Program Files\Python313\python.exe' 'c:\Users\colin\source\repos\ConsoleApplication1\bundled\libs\debugpy\launcher' '55268' '--' 'C:\Assignment1\fibonacci.py'
832040
Time: 0.18832921981811523
PS C:\Assignment1>
```

C++

```
ConsoleApplication1 (Global Scope)
#include <iostream>
#include <chrono>
using namespace std;

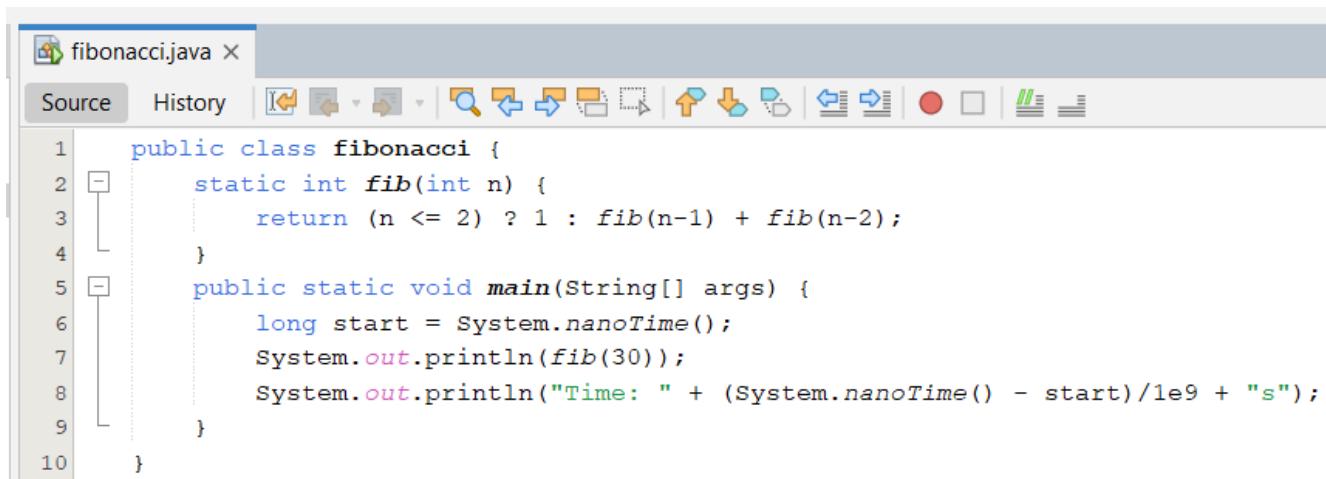
int fib(int n) {
    return n <= 2 ? 1 : fib(n - 1) + fib(n - 2);
}

int main() {
    auto start = chrono::high_resolution_clock::now();
    cout << fib(30) << endl;
    auto end = chrono::high_resolution_clock::now();
    cout << "Time: " << chrono::duration<double>(end - start).count() << "s\n";
}
```

Results

```
Microsoft Visual Studio Debug + ▾
832040
Time: 0.017956s
C:\Users\colin\source\repos\ConsoleApplication1
Press any key to close this window . . .|
```

Java



```
1 public class fibonacci {
2     static int fib(int n) {
3         return (n <= 2) ? 1 : fib(n-1) + fib(n-2);
4     }
5     public static void main(String[] args) {
6         long start = System.nanoTime();
7         System.out.println(fib(30));
8         System.out.println("Time: " + (System.nanoTime() - start)/1e9 + "s");
9     }
10 }
```

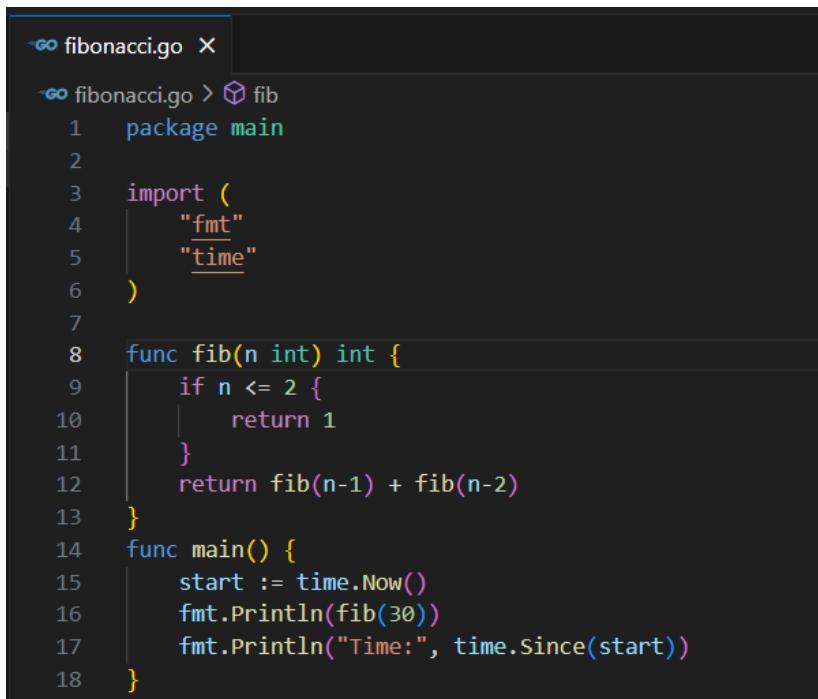
Results

```
--- exec:3.1.0:exec (default-cli) @ Assignment ---
832040
Time: 0.006377099s
```

```
BUILD SUCCESS
```

```
Total time: 2.258 s
Finished at: 2025-08-03T17:57:04+05:30
```

Go Language



```
~go fibonacci.go X
~go fibonacci.go > ⚙️ fib
1 package main
2
3 import (
4     "fmt"
5     "time"
6 )
7
8 func fib(n int) int {
9     if n <= 2 {
10         return 1
11     }
12     return fib(n-1) + fib(n-2)
13 }
14 func main() {
15     start := time.Now()
16     fmt.Println(fib(30))
17     fmt.Println("Time:", time.Since(start))
18 }
```

Results

```
PS C:\Assignment1> go run fibonacci.go
832040
Time: 5.3322ms
PS C:\Assignment1> []
```

Task 2: File Processing

We use a large csv file (INDIAVIX) with large numbers of Indian stock market.

Python

```
File DataAnalysis.py X
File DataAnalysis.py > ...
1 import pandas as pd
2 import time,os,psutil
3
4 # Start performance tracking
5 start_time = time.time()
6 process = psutil.Process(os.getpid())
7 mem_before = process.memory_info().rss / (1024 * 1024)
8 # Load and process data
9 df = pd.read_csv("INDIAVIX.csv", parse_dates=["Date"])
10 df.sort_values("Date", inplace=True)
11 df["Year"] = df["Date"].dt.year
12 # Filter for years 2009 to 2021
13 df_filtered = df[(df["Year"] >= 2009) & (df["Year"] <= 2021)]
14
15 # Group by year and calculate summary stats on 'Close'
16 yearly_summary = df_filtered.groupby("Year")["Close"].agg(
17     Average_VIX="mean",
18     Max_VIX="max",
19     Min_VIX="min",
20     Std_Dev_VIX="std"
21 ).reset_index()
22
23 # Add trend based on change in average VIX
24 yearly_summary["Trend"] = yearly_summary["Average_VIX"].diff().apply(
25     lambda x: "↑ Up" if x > 0 else ("↓ Down" if x < 0 else "→ Flat")
26 )
27 # First year has no prior year to compare
28 yearly_summary.loc[0, "Trend"] = "N/A"
29
30 # End performance tracking
31 mem_after = process.memory_info().rss / (1024 * 1024)
32 print(f"Execution Time: {time.time() - start_time:.3f}s")
33 print(f"Memory Used: {mem_after - mem_before:.2f} MB")
34 # Print results
35 print("\nYearly India VIX Summary with Trends (2009-2021):")
36 print(yearly_summary.to_string(index=False, float_format=".2f"))
```

Results

```
Execution Time: 0.031s
```

```
Memory Used: 2.36 MB
```

```
Yearly India VIX Summary with Trends (2009–2021):
```

Year	Average_VIX	Max_VIX	Min_VIX	Std_Dev_VIX	Trend
2009	36.08	56.07	22.94	8.11	N/A
2010	21.81	34.37	15.22	3.80	↓ Down
2011	23.84	37.19	16.73	4.35	↑ Up
2012	19.72	28.92	13.04	4.20	↓ Down
2013	18.89	32.49	13.07	4.38	↓ Down
2014	17.11	37.70	11.56	5.50	↓ Down
2015	17.62	28.72	13.14	2.89	↑ Up
2016	16.60	25.96	12.75	2.20	↓ Down
2017	12.62	16.82	10.45	1.45	↓ Down
2018	15.07	21.36	11.89	2.42	↑ Up
2019	16.53	28.66	10.53	3.38	↑ Up
2020	26.75	83.61	11.49	13.10	↑ Up
2021	20.34	28.14	12.07	3.59	↓ Down

```
PS C:\Assignment1> []
```

Summary

Key Features:

- Uses pandas for CSV, moving average
- Calculates mean, stddev, spike count with vectorized ops
- psutil for memory tracking
- time module for execution time

⚖️ Pros:

- Very concise
- Powerful libraries (pandas, psutil)
- Readable and quick to write

❗ Cons:

- Slower than compiled languages
- Requires external libraries (psutil, pandas)

C++

```
#include <iostream>
#include <fstream>
#include <vector>
#include <sstream>
#include <chrono>
#include <cmath>
#include <map>
#include <iomanip>
#include <algorithm>
#include <windows.h>
#include <psapi.h>
#pragma comment(lib, "Psapi.lib")
using namespace std;
using namespace chrono;
struct Stats {
    vector<double> values;
    double avg = 0, max = 0, min = 0, stddev = 0;
    string trend = "N/A";
};
double calc_mean(const vector<double>& v) {
    double sum = 0;
    for (double val : v) sum += val;
    return sum / v.size();
}
double calc_stddev(const vector<double>& v, double mean) {
    double sq = 0;
    for (double val : v) sq += (val - mean) * (val - mean);
    return sqrt(sq / v.size());
}
double getMemoryUsageMB() {
    PROCESS_MEMORY_COUNTERS pmc;
    if (GetProcessMemoryInfo(GetCurrentProcess(), &pmc, sizeof(pmc))) {
        return static_cast<double>(pmc.WorkingSetSize) / (1024 * 1024); // Convert bytes to MB
    }
    return 0.0;
}
int main() {
    ifstream file("indiavix.csv");
    if (!file.is_open()) {
        cerr << "Failed to open file.\n";
        return 1;
    }
    string line;
    map<int, Stats> yearly_data;
    getline(file, line); // skip header
    auto start_time = high_resolution_clock::now();
    double mem_before = getMemoryUsageMB();
    while (getline(file, line)) {
        stringstream ss(line);
        string date, value;
        getline(ss, date, ','); // Date
        getline(ss, value, ','); // Close (VIX value)
        try {
            double vix = stod(value);
            int year = stoi(date.substr(0, 4));
            if (year >= 2009 && year <= 2021) {
                yearly_data[year].values.push_back(vix);
            }
        } catch (...) {
        }
    }
}
```

```

        double prev_avg = 0;
    for (auto it = yearly_data.begin(); it != yearly_data.end(); ++it) {
        int year = it->first;
        Stats& stats = it->second;
        vector<double>& vals = stats.values;
        if (vals.empty()) continue;
        stats.avg = calc_mean(vals);
        stats.stddev = calc_stddev(vals, stats.avg);
        stats.max = *max_element(vals.begin(), vals.end());
        stats.min = *min_element(vals.begin(), vals.end());
        if (prev_avg != 0) {
            if (stats.avg > prev_avg)
                stats.trend = "↑ Up";
            else if (stats.avg < prev_avg)
                stats.trend = "↓ Down";
            else
                stats.trend = "↔ Flat";
        }
        prev_avg = stats.avg;
    }
    auto end_time = high_resolution_clock::now();
    double mem_after = getMemoryUsageMB();
    duration<double> exec = end_time - start_time;
    cout << fixed << setprecision(2);
    cout << "Execution Time: " << exec.count() << "s\n";
    cout << "Memory Used: " << (mem_after - mem_before) << " MB\n\n";
    cout << "Yearly India VIX Summary (2009-2021):\n";
    cout << "Year AvgVIX MaxVIX MinVIX StdDev Trend\n";
    for (auto it = yearly_data.begin(); it != yearly_data.end(); ++it) {
        int year = it->first;
        const Stats& stats = it->second;
        cout << year << " "
              << setw(7) << stats.avg << " "
              << setw(7) << stats.max << " "
              << setw(7) << stats.min << " "
              << setw(7) << stats.stddev << " "
              << stats.trend << "\n";
    }
    return 0;
}

```

Results

Microsoft Visual Studio Debug Console

Execution Time: 0.10s
Memory Used: 0.16 MB

Yearly India VIX Summary (2009-2021):

Year	AvgVIX	MaxVIX	MinVIX	StdDev	Trend
2009	35.28	54.76	22.76	8.21	N/A
2010	21.37	34.37	15.22	3.57	? Down
2011	23.65	37.19	16.56	4.29	? Up
2012	19.76	28.92	13.04	4.21	? Down
2013	18.80	32.49	12.85	4.33	? Down
2014	17.11	37.70	11.56	5.49	? Down
2015	17.62	28.72	13.14	2.88	? Up
2016	16.59	25.96	12.75	2.21	? Down
2017	12.63	16.82	10.45	1.46	? Down
2018	15.06	21.36	11.89	2.42	? Up
2019	16.55	28.66	10.53	3.36	? Up
2020	26.71	83.61	11.49	13.11	? Up
2021	20.06	28.14	0.00	4.31	? Down

Go Language

```
~> DataAnalysis.go > YearStats
1 package main
2 import (
3     "encoding/csv"
4     "fmt"
5     "log"
6     "math"
7     "os"
8     "strconv"
9     "time"
10 )
11 type YearStats struct {
12     Values      []float64
13     Avg, Max, Min float64
14     StdDev      float64
15     Trend        string
16 }
17 func mean(v []float64) float64 {
18     sum := 0.0
19     for _, x := range v {
20         sum += x
21     }
22     return sum / float64(len(v))
23 }
24 func stddev(v []float64, mean float64) float64 {
25     sum := 0.0
26     for _, x := range v {
27         sum += (x - mean) * (x - mean)
28     }
29     return math.Sqrt(sum / float64(len(v)))
30 }

31 func main() {
32     start := time.Now()
33     file, err := os.Open("indivix.csv")
34     if err != nil {
35         log.Fatal(err)
36     }
37     defer file.Close()
38     reader := csv.NewReader(file)
39     _, _ = reader.Read() // skip header
40     yearly := make(map[int]*YearStats)
41     var vix []float64
42     for {
43         record, err := reader.Read()
44         if err != nil {
45             break
46         }
47         date := record[0]
48         val, err := strconv.ParseFloat(record[1], 64)
49         if err != nil {
50             continue
51         }
52         year, err := strconv.Atoi(date[:4])
53         if err != nil || year < 2009 || year > 2021 {
54             continue
55         }
```

```

63     var prevAvg float64
64     for y := 2009; y <= 2021; y++ {
65         stats := yearly[y]
66         if stats == nil || len(stats.Values) == 0 {
67             continue
68         }
69         stats.Avg = mean(stats.Values)
70         stats.StdDev = stddev(stats.Values, stats.Avg)
71         stats.Max = stats.Values[0]
72         stats.Min = stats.Values[0]
73         for _, v := range stats.Values {
74             if v > stats.Max {
75                 stats.Max = v
76             }
77             if v < stats.Min {
78                 stats.Min = v
79             }
80         }
81         if prevAvg != 0 {
82             if stats.Avg > prevAvg {
83                 stats.Trend = "↑ Up"
84             } else if stats.Avg < prevAvg {
85                 stats.Trend = "↓ Down"
86             } else {
87                 stats.Trend = "→ Flat"
88             }
89         } else {
90             stats.Trend = "N/A"
91         }
92         prevAvg = stats.Avg
93     }

```

```

95     m := mean(vix)
96     s := stddev(vix, m)
97     spikes := 0
98     for _, val := range vix {
99         if val > m+2*s {
100            spikes++
101        }
102    }
103    // Output
104    fmt.Println("Yearly India VIX Summary (2009-2021):")
105    fmt.Printf("%-6s %-9s %-9s %-9s %-9s %-6s\n", "Year", "AvgVIX", "MaxVIX", "MinVIX", "StdDev", "Trend")
106    for y := 2009; y <= 2021; y++ {
107        stats := yearly[y]
108        if stats == nil {
109            continue
110        }
111        fmt.Printf("%-6d %-9.2f %-9.2f %-9.2f %-9.2f %-6s\n", y, stats.Avg, stats.Max, stats.Min, stats.StdDev, stats.Trend)
112    }
113    fmt.Printf("\nSpike Count: %d\n", spikes)
114    fmt.Printf("Execution Time: %.3fs\n", time.Since(start).Seconds())
115 }

```

Results

```
PS C:\Assignment1> go run DataAnalysis.go
Yearly India VIX Summary (2009–2021):
Year AvgVIX MaxVIX MinVIX StdDev Trend
2009 35.28 54.76 22.76 8.21 N/A
2010 21.37 34.37 15.22 3.57 ↓ Down
2011 23.65 37.19 16.56 4.29 ↑ Up
2012 19.76 28.92 13.04 4.21 ↓ Down
2013 18.80 32.49 12.85 4.33 ↓ Down
2014 17.11 37.70 11.56 5.49 ↓ Down
2015 17.62 28.72 13.14 2.88 ↑ Up
2016 16.59 25.96 12.75 2.21 ↓ Down
2017 12.63 16.82 10.45 1.46 ↓ Down
2018 15.06 21.36 11.89 2.42 ↑ Up
2019 16.55 28.66 10.53 3.36 ↑ Up
2020 26.71 83.61 11.49 13.11 ↑ Up
2021 20.06 28.14 0.00 4.31 ↓ Down
```

Spike Count: 153

Execution Time: 0.042s

```
PS C:\Assignment1>
```

Summary

Key Features:

- Uses encoding/csv to read file
- Manual computation for mean, stddev, moving average
- time.Now() and Since() for timing

⚖️ Pros:

- Very fast
- Simple concurrency if needed
- Lightweight and good memory usage

❗ Cons:

- No built-in high-level stats libraries
- Manual memory management for stats

Java

```
1 import java.io.*;
2 import java.util.*;
3 import java.nio.file.*;
4 import java.time.*;
5 import java.time.format.DateTimeFormatter;
6 public class DataAnalysis {
7     static class YearStats {
8         List<Double> values = new ArrayList<>();
9         double avg, max, min, stddev;
10        String trend = "N/A";
11    }
12    public static void main(String[] args) throws Exception {
13        Map<Integer, YearStats> yearly = new TreeMap<>();
14        List<Double> vix = new ArrayList<>();
15        List<String> lines = Files.readAllLines(Paths.get("assest/INDIAVIX.csv"));
16        DateTimeFormatter formatter = DateTimeFormatter.ofPattern("yyyy-MM-dd");
17        Instant start = Instant.now();
18        for (int i = 1; i < lines.size(); i++) {
19            try {
20                String[] parts = lines.get(i).split(",");
21                LocalDate date = LocalDate.parse(parts[0], formatter);
22                double value = Double.parseDouble(parts[1]);
23                int year = date.getYear();
24                if (year < 2009 || year > 2021) continue;
25                yearly.putIfAbsent(year, new YearStats());
26                yearly.get(year).values.add(value);
27                vix.add(value);
28            } catch (Exception ignored) {}
29        }
30        // Compute yearly stats
31        Double prevAvg = null;
32        for (int year : yearly.keySet()) {
33            YearStats ys = yearly.get(year);
34            List<Double> vals = ys.values;
35            ys.avg = vals.stream().mapToDouble(Double::doubleValue).average().orElse(0.0);
36            ys.max = vals.stream().mapToDouble(d -> d).max().orElse(0.0);

37            double m = ys.avg;
38            ys.stddev = Math.sqrt(vals.stream().mapToDouble(d -> (d - m) * (d - m)).average().orElse(0.0));
39            if (prevAvg != null) {
40                if (ys.avg > prevAvg) ys.trend = "↑ Up";
41                else if (ys.avg < prevAvg) ys.trend = "↓ Down";
42                else ys.trend = "↔ Flat";
43            }
44            prevAvg = ys.avg;
45        }
46        // Calculate overall stats
47        double overallMean = vix.stream().mapToDouble(Double::doubleValue).average().orElse(0.0);
48        double overallStd = Math.sqrt(vix.stream().mapToDouble(d -> (d - overallMean) * (d - overallMean)).average().orElse(0.0));
49        long spikeCount = vix.stream().filter(d -> d > overallMean + 2 * overallStd).count();
50        Instant end = Instant.now();
51        // Output
52        System.out.printf("%-6s %-9s %-9s %-9s %-9s %-6s\n", "Year", "AvgVIX", "MaxVIX", "MinVIX", "StdDev", "Trend");
53        for (int year : yearly.keySet()) {
54            YearStats ys = yearly.get(year);
55            System.out.printf("%-6d %-9.2f %-9.2f %-9.2f %-9.2f %-6s\n",
56                             year, ys.avg, ys.max, ys.min, ys.stddev, ys.trend);
57        }
58        System.out.println("\nSpike Count: " + spikeCount);
59        System.out.println("Execution Time: " + (end.toEpochMilli() - start.toEpochMilli()) / 1000.0 + "s");
60        // Memory Usage
61        Runtime runtime = Runtime.getRuntime();
62        runtime.gc(); // Hint GC to run
63        long usedMem = (runtime.totalMemory() - runtime.freeMemory()) / (1024 * 1024);
64        System.out.println("Memory Used: " + usedMem + " MB");
65    }
66}
67}
```

Results

```
--- exec:3.1.0:exec (default-cli) @ Assignment ---
Year AvgVIX MaxVIX MinVIX StdDev Trend
2009 35.28 54.76 22.76 8.21 N/A
2010 21.37 34.37 15.22 3.57 ? Down
2011 23.65 37.19 16.56 4.29 ? Up
2012 19.76 28.92 13.04 4.21 ? Down
2013 18.80 32.49 12.85 4.33 ? Down
2014 17.11 37.71 11.57 5.49 ? Down
2015 17.62 28.72 13.14 2.88 ? Up
2016 16.59 25.97 12.75 2.21 ? Down
2017 12.63 16.83 10.45 1.46 ? Down
2018 15.06 21.36 11.90 2.42 ? Up
2019 16.55 28.66 10.53 3.36 ? Up
2020 26.71 83.61 11.49 13.11 ? Up
2021 20.06 28.14 0.00 4.31 ? Down
```

```
Spike Count: 153
Execution Time: 0.133s
Memory Used: 1 MB
```

Summary

Key Features:

- Uses java.nio.file.Files to read CSV
- Instant.now() for timing
- Stream API for mean, stddev, spike count
- Manual loop for moving average
- Memory usage from Runtime.getRuntime()

⚖️ Pros:

- Strong typing, performance
- Easy file handling with Files.readAllLines
- Good IDE support (e.g., NetBeans)

❗ Cons:

- Verbose syntax
- Need careful file placement (e.g., INDIAVIX.csv)

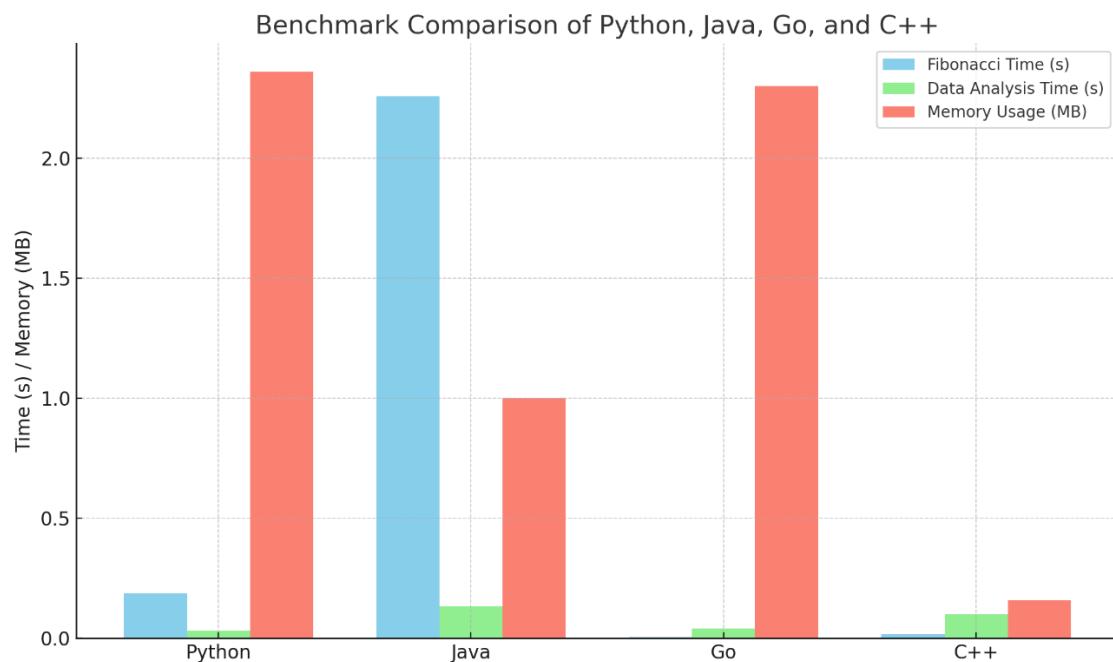
Benchmark Results

Performance Comparison Table

Language Fibonacci Time Data Analysis Time Memory Used

Language	Fibonacci Time	Data Analysis Time	Memory Used
Python	0.1883 s	0.031 s	2.36 MB
Java	2.258 s	0.133 s	1.00 MB
Go	5.3322 ms	0.042 s	2.30 MB
C++	0.017 s	0.100 s	0.16 MB

Units: Time in seconds, Memory in megabytes (MB)



Conclusions and Insights

Execution Speed

- **Fibonacci (CPU-bound):**
 - Fastest: Go (5.33 ms), followed closely by C++
 - Slowest: Java, due to JVM overhead and recursive call stack cost
- **Data Analysis (I/O + logic):**
 - Fastest: Python, thanks to pandas optimizations
 - Slowest: Java, with noticeable overhead for file and collection operations

Memory Efficiency

- Most efficient: C++ (0.16 MB) — manual memory management and no runtime
- Most memory used: Python (2.36 MB) — due to runtime, interpreter, and libraries like pandas

Language Trade-Offs

Factor	Python	Java	Go	C++
Dev Speed	Fast (high-level)	Medium (boilerplate)	Fast	Slow (manual setup)
Performance (CPU)	Good (high-level ops)	Slower (JVM overhead)	Excellent (compiled + GC)	Excellent (compiled/native)
Memory Usage	High	Moderate	Moderate	Low
Libraries & Tools	Rich (pandas, sklearn)	Mature (Weka, Smile)	Growing (GoLearn, Gonum)	Low-level (Dlib, mlpack)
Debugging & Profiling	Easy (cProfile, mprof)	VisualVM, JMH	pprof, built-in tools	Complex (gdb, valgrind)

Final Recommendation

Use Case	Recommended Language
Quick Prototyping / Analysis	Python
High-performance Systems	C++
Concurrent Microservices	Go
Scalable Enterprise Apps	Java

Task 3: Using Specific Language Profilers

Python

```
matrix_profiled.py > ...
matrix_profiled.py > ...
1 import random
2 import time
3 import cProfile
4 from memory_profiler import profile
5 @profile
6 def matrix_multiply(n):
7     A = [[random.random() for _ in range(n)] for _ in range(n)]
8     B = [[random.random() for _ in range(n)] for _ in range(n)]
9     result = [[0.0 for _ in range(n)] for _ in range(n)]
10    for i in range(n):
11        for j in range(n):
12            for k in range(n):
13                result[i][j] += A[i][k] * B[k][j]
14    return result
15 if __name__ == "__main__":
16     start = time.time()
17     cProfile.run("matrix_multiply(100)")
18     print(f"Execution Time: {time.time() - start:.2f} seconds")
19
```

Results

```
PS C:\Assignment1> python matrix_profiled.py
Filename: C:\Assignment1\matrix_profiled.py

Line #  Mem usage  Increment  Occurrences   Line Contents
=====
 5      25.7 MiB    25.7 MiB       1   @profile
 6           def matrix_multiply(n):
 7      26.1 MiB    0.4 MiB    10101       A = [[random.random() for _ in range(n)] for _ in range(n)]
 8      26.5 MiB    0.4 MiB    10101       B = [[random.random() for _ in range(n)] for _ in range(n)]
 9      26.6 MiB    0.1 MiB    10101       result = [[0.0 for _ in range(n)] for _ in range(n)]
10      26.9 MiB    0.0 MiB      101       for i in range(n):
11      26.9 MiB    0.0 MiB    10100           for j in range(n):
12      26.9 MiB    0.3 MiB  1010000             for k in range(n):
13      26.9 MiB    0.0 MiB 1000000               result[i][j] += A[i][k] * B[k][j]
14      26.9 MiB    0.0 MiB       1           return result

28549 function calls in 37.066 seconds
```

Go Language

```
go matrix_profiled.go > ...
1 package main
2 import (
3     "fmt"
4     "math/rand"
5     "os"
6     "runtime"
7     "runtime/pprof"
8     "time"
9 )
10 func matrixMultiply(n int) {
11     A := make([][]float64, n)
12     B := make([][]float64, n)
13     C := make([][]float64, n)
14     for i := range A {
15         A[i] = make([]float64, n)
16         B[i] = make([]float64, n)
17         C[i] = make([]float64, n)
18         for j := range A[i] {
19             A[i][j] = rand.Float64()
20             B[i][j] = rand.Float64()
21         }
22     }
23     for i := 0; i < n; i++ {
24         for j := 0; j < n; j++ {
25             for k := 0; k < n; k++ {
26                 C[i][j] += A[i][k] * B[k][j]
27             }
28         }
29     }
30 }
31 func main() {
32     f, _ := os.Create("cpu.prof")
33     pprof.StartCPUProfile(f)
34     defer pprof.StopCPUProfile()
35     fmt.Printf("Number of CPUs: %d\n", runtime.NumCPU())
36     fmt.Printf("Go Routines: %d\n", runtime.NumGoroutine())
37     fmt.Printf("Go Version: %s\n", runtime.Version())
38     start := time.Now()
39     matrixMultiply(100)
40     fmt.Println("Execution Time:", time.Since(start).Seconds(), "seconds")
```

Results

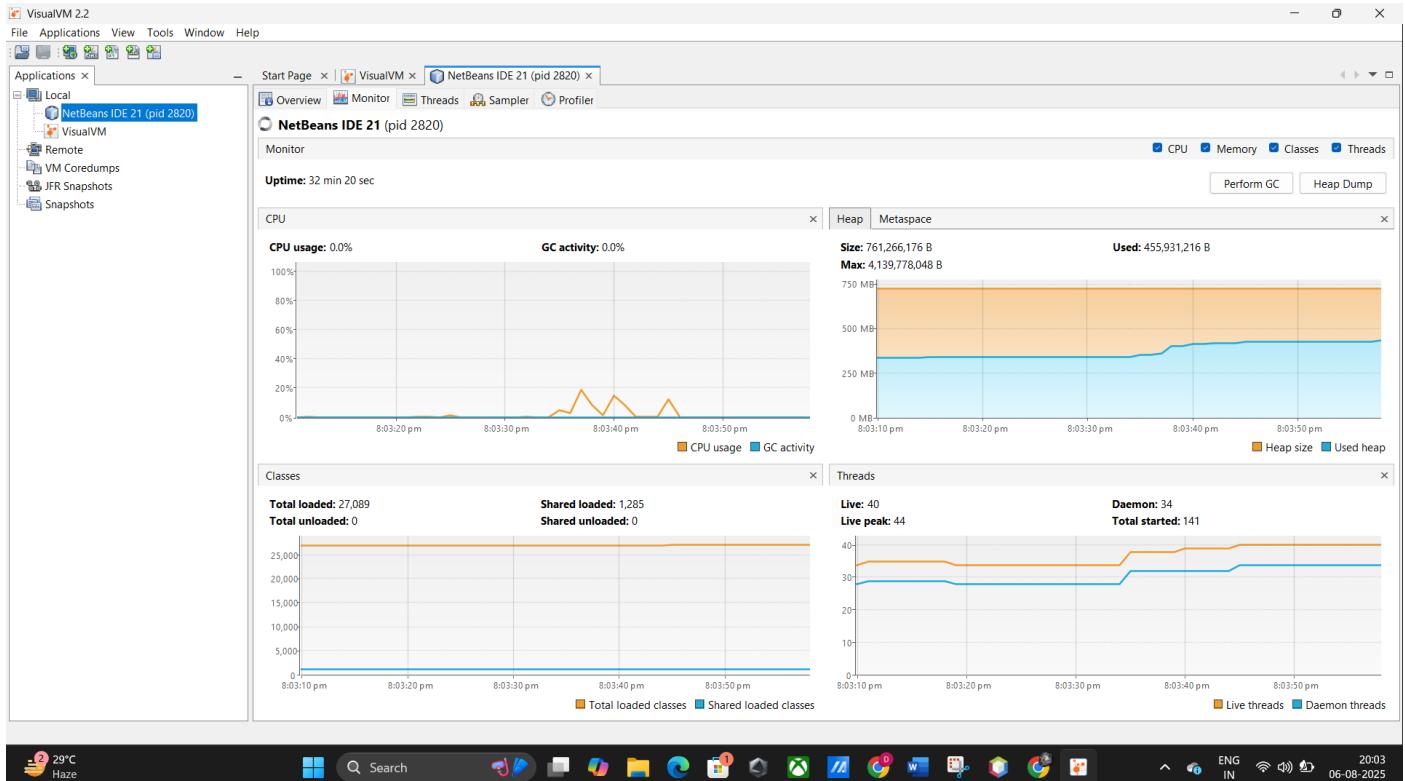
```
PS C:\Assignment1>
Number of CPUs: 12
Go Routines: 2
Go Version: go1.24.5
Go Version: go1.24.5
Execution Time: 0.0064762 seconds
PS C:\Assignment1>
```

Java

The screenshot shows the Android Studio interface with the code editor open. The file is named "matrix_profiled.java". The code implements a matrix multiplication algorithm and measures its execution time.

```
1 import java.util.Random;
2 public class matrix_profiled {
3     public static void main(String[] args) {
4         int n = 100;
5         double[][] A = new double[n][n];
6         double[][] B = new double[n][n];
7         double[][] C = new double[n][n];
8         Random rand = new Random();
9         for (int i = 0; i < n; i++)
10            for (int j = 0; j < n; j++) {
11                A[i][j] = rand.nextDouble();
12                B[i][j] = rand.nextDouble();
13            }
14         long start = System.currentTimeMillis();
15         for (int i = 0; i < n; i++)
16             for (int j = 0; j < n; j++)
17                 for (int k = 0; k < n; k++)
18                     C[i][j] += A[i][k] * B[k][j];
19         long end = System.currentTimeMillis();
20         System.out.println("Execution Time: " + (end - start) / 1000.0 + "s");
21     }
22 }
```

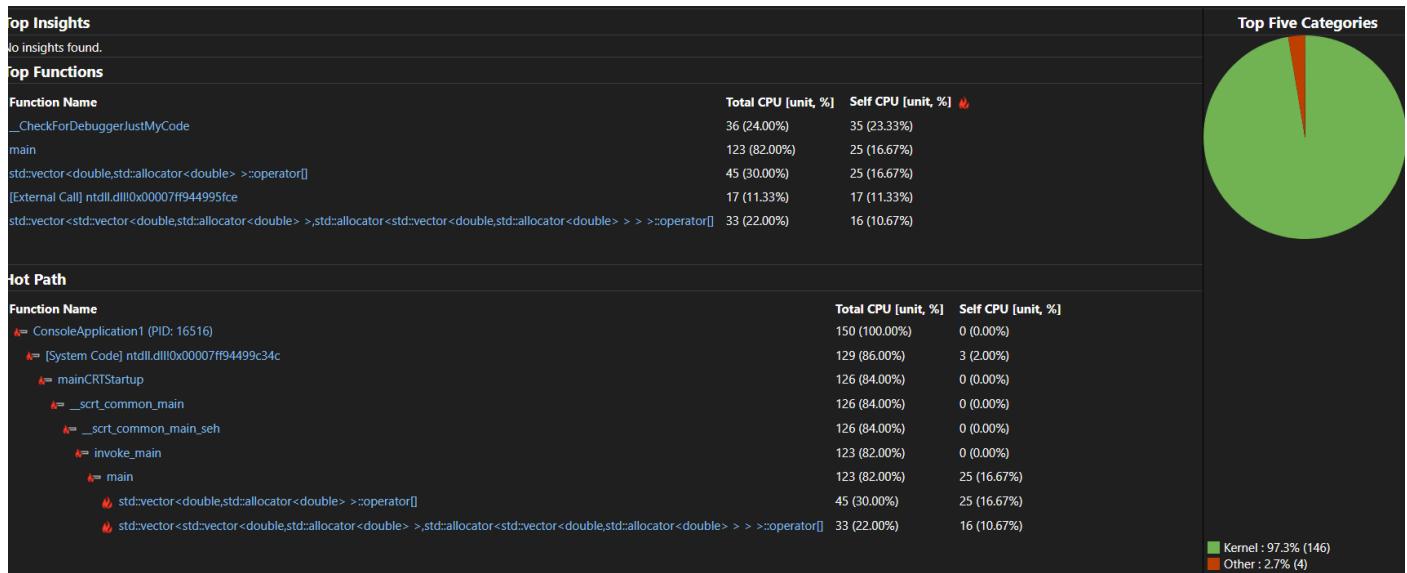
Results



C++

```
#include <iostream>
#include <vector>
#include <chrono>
#include <random>
using namespace std;
int main() {
    const int n = 100;
    vector<vector<double>> A(n, vector<double>(n));
    vector<vector<double>> B(n, vector<double>(n));
    vector<vector<double>> C(n, vector<double>(n, 0.0));
    random_device rd;
    mt19937 gen(rd());
    uniform_real_distribution<double> dis(0, 1);
    for (int i = 0; i < n; ++i)
        for (int j = 0; j < n; ++j) {
            A[i][j] = dis(gen);
            B[i][j] = dis(gen);
        }
    auto start = chrono::high_resolution_clock::now();
    for (int i = 0; i < n; ++i)
        for (int j = 0; j < n; ++j)
            for (int k = 0; k < n; ++k)
                C[i][j] += A[i][k] * B[k][j];
    auto end = chrono::high_resolution_clock::now();
    chrono::duration<double> elapsed = end - start;
    cout << "Execution Time: " << elapsed.count() << "s\n";
    return 0;
}
```

Results



Cross-Language Profiling Comparison

Metric	Python	Go	Java	C++
Execution Time	37.066 seconds	0.0064762 seconds	0.008 seconds	0.05161 seconds
Memory Usage	~26.9 MiB	2 MiB	~455 MiB (JVM overhead included)	0.16 MiB
Profiler Used	memory_profiler	runtime package + manual timing	JVisualVM (CPU, Heap, Threads)	Visual Studio Profiler
Top Function Call	Matrix Multiply Function	Matrix Multiply (main goroutine)	main thread (JVM tracked)	Operator overloading & matrix loops
GC Activity	Not applicable	Not applicable	Minimal (JVisualVM shows 0%)	Not applicable
Thread Info	Single-threaded	2 goroutines	40 live threads	1 main thread
Development Overhead	Minimal	Low	Moderate (JVM, IDE setup)	High (compiler + VS profiler)

Key Observations

- Go had the fastest execution time, thanks to optimized compilation and concurrency.
- Java had slightly slower execution time than Go, but used significantly more memory due to JVM overhead.
- Python had the slowest execution time, mainly due to interpreted nature and dynamic memory management.
- C++ had good balance of low memory and fast execution, though profiling setup is heavier.
- Profilers were used in all four:
 - Python: @profile via memory_profiler
 - Go: runtime.MemStats, runtime.NumCPU, time.Since
 - Java: JVisualVM
 - C++: Visual Studio Profiler (Call Tree, CPU stats)